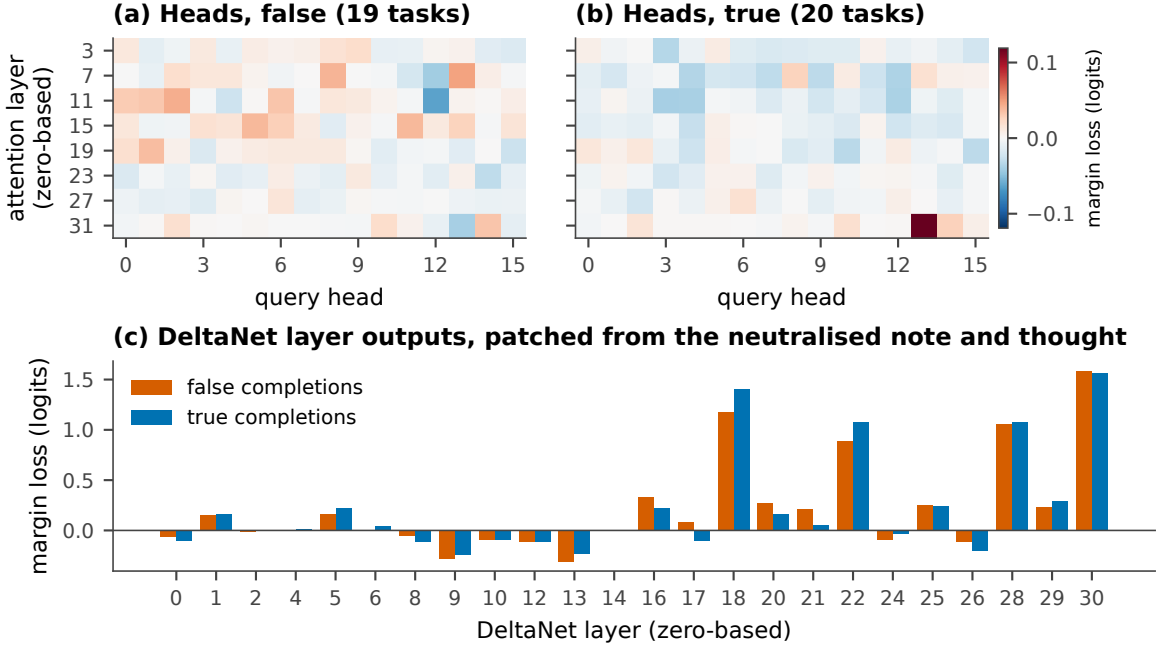
(a) Heads, false (19 tasks)
(b) Heads, true (20 tasks)
attention layer (zero-based)
3
7
11
15
19
23
27
31
0
3
6
9
12
15
query head
0.1
0.0
−0.1
margin loss (logits)
(c) DeltaNet layer outputs, patched from the neutralised note and thought
false completions
true completions
margin loss (logits)
1.5
1.0
0.5
0.0
0
1
2
4
5
6
8
9
10
12
13
14
16
17
18
20
21
22
24
25
26
28
29
30
DeltaNet layer (zero-based)